



# Network Failure Prediction on CNFs 5GC with Linux eBPF

UT-NakaoLab-AI

Takeru Hakii


# agenda

- Introduction
- Dataset description
- Task1
- Task1 method
- Task1 result
- Task2
- Task2 method
- Task2 result
- Feature analysis
- Visualization
- Conclusion
- Future works

# summary

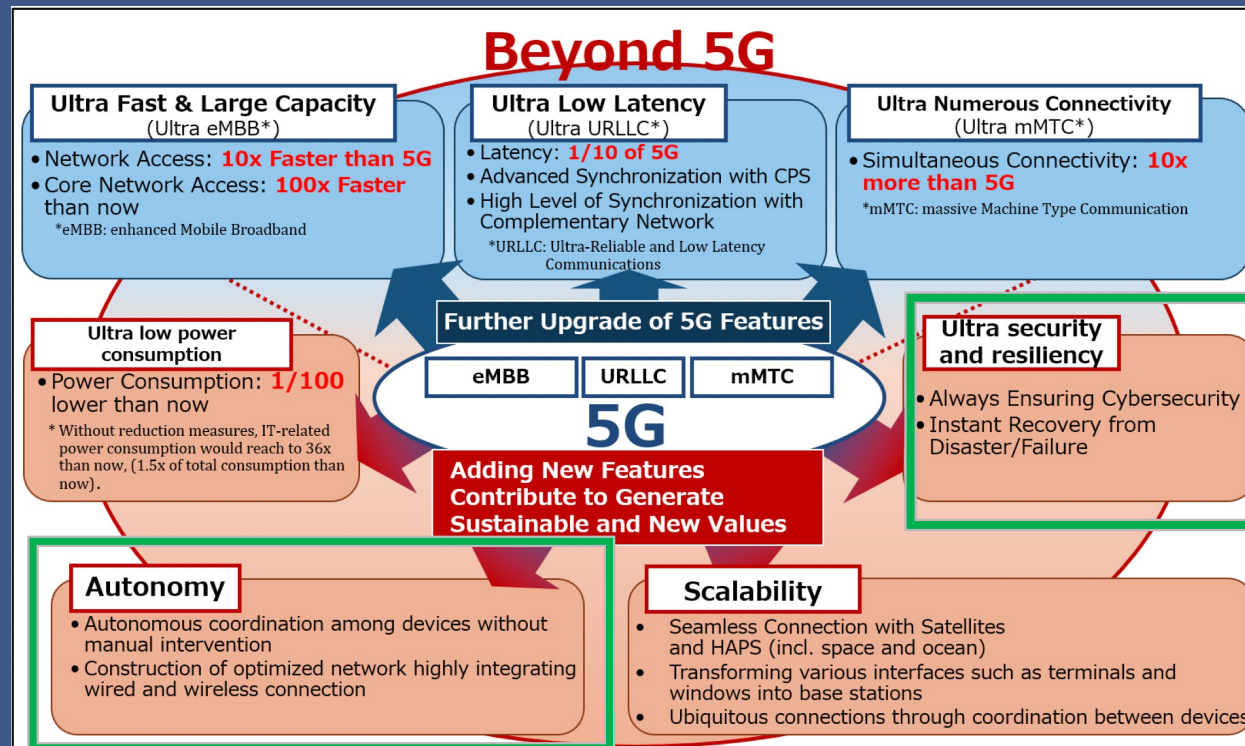
- In task1, my model achieved f1 score=0.92 with detection time=140s
- In task2, my model successfully reduced the number of features from 3325 to 329
- My best solution resulted in 329 features, detection time=140s, f1 score=0.97

# Introduction




By leveraging flexible containerized network functions (NF)  
Scalable networks can be rapidly deployed

# Introduction



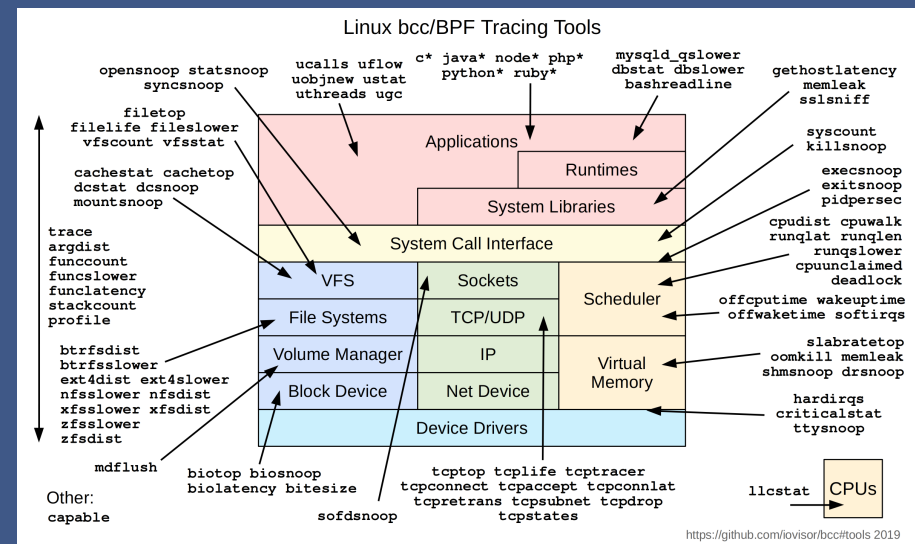
Need a system that monitors the network and automatically detects and restores failures.

# Introduction

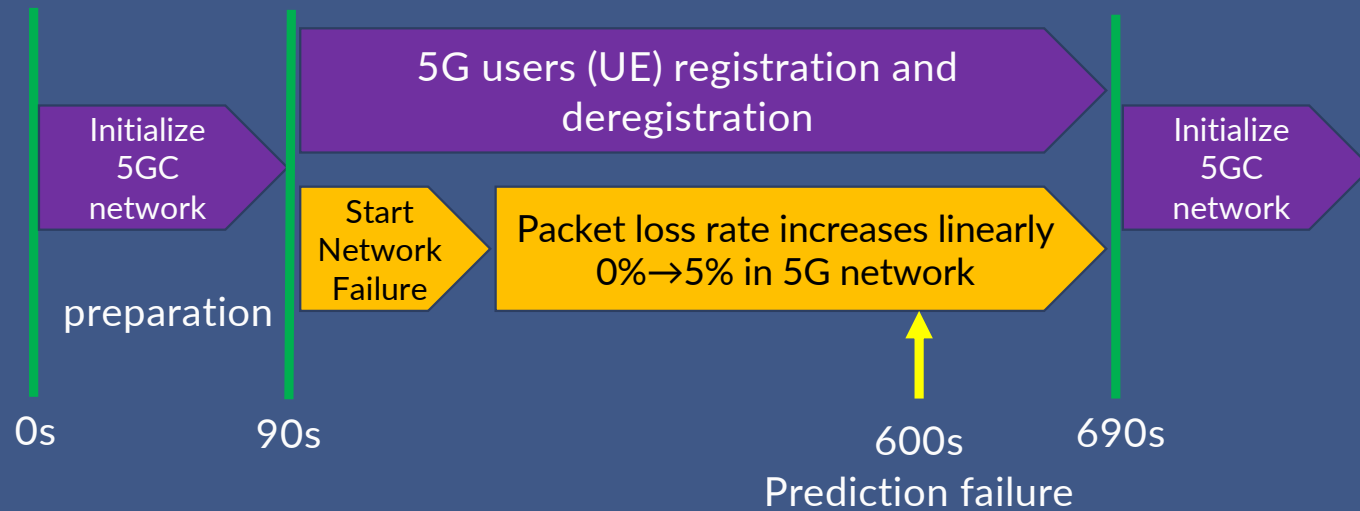


Only log data obtained from containers  
difficult to predict failures in advance

Using eBPF (extended Berkley Packet Filter),  
we get more detailed data





# Dataset description



- One cycle is 690 seconds
- Logging every 10 seconds
- The first 90 seconds is a preparation period, initializing the network
- The subsequent 600 seconds (10 minutes) cause a failure, increasing the packet loss rate linearly

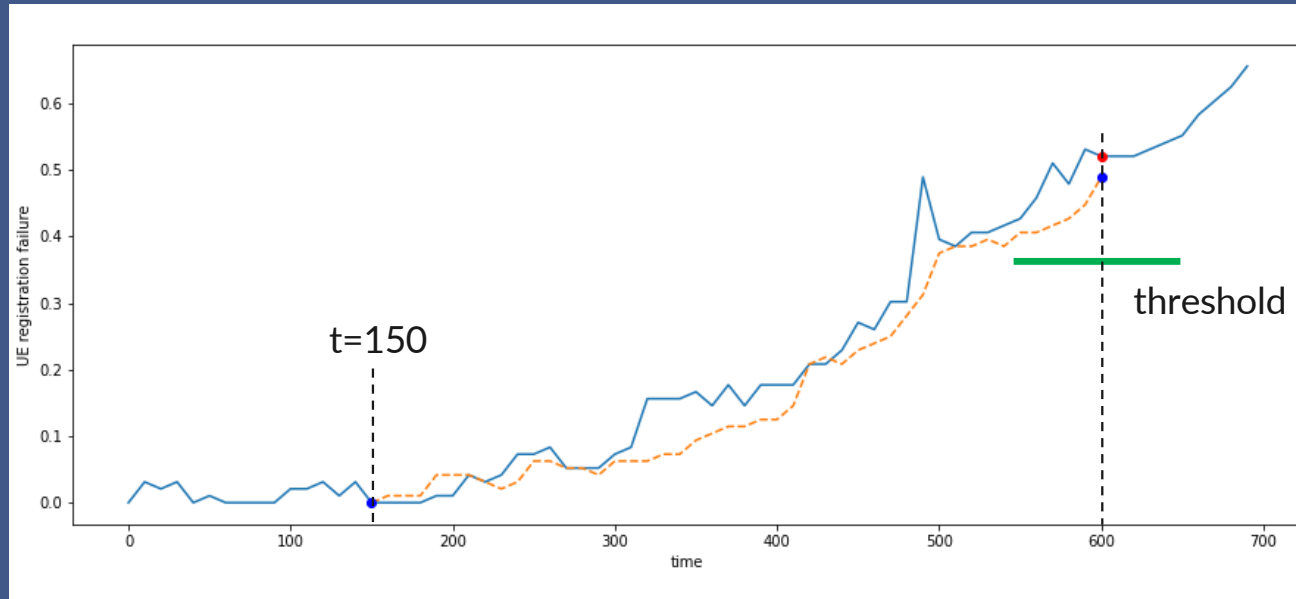
# Dataset description

- Train : 600 cycles
- Test : 300 cycles
- Metrics : 3326





# Task1



Predicts the number of UE registration failures at 600 seconds from a certain point  $t$ , and determines if a failure has occurred according to a threshold value.

Make  $t$  as small as possible so that the F1 score exceeds 0.9.


# Method(Task1)

- Use all features (all) and remove features that are 0 for all periods (removed)
- Feature size
  - all: 3325
  - removed: 1723
- Scaling of both training and test data using minimum and maximum values of training data
- Threshold = 0.3 (28.8)
- timesteps: the input sequence size
- delay: interval between input and output

- Model: LSTM
- Loss function: MSE
- Optimizer: adam
- Epochs: 15
- Training data: 500 cycles
- Validation data: 100 cycles

| Detection time $t$ | timesteps | delay | Data size | Loss input |
|--------------------|-----------|-------|-----------|------------|
| 250                | 15        | 35    | 10000     | 15         |
| 200                | 15        | 40    | 7500      | 10         |
| 150                | 15        | 45    | 5000      | 5          |
| 140                | 14        | 46    | 5000      | 4          |
| 130                | 13        | 47    | 5000      | 3          |
| 120                | 12        | 48    | 5000      | 2          |

# Removed model has high performance



Successful prediction at detection time 150s by using all features  
Successful prediction at detection time 140s by removing 0 metrics  
The removed model has higher performance than the all model.

# Task2

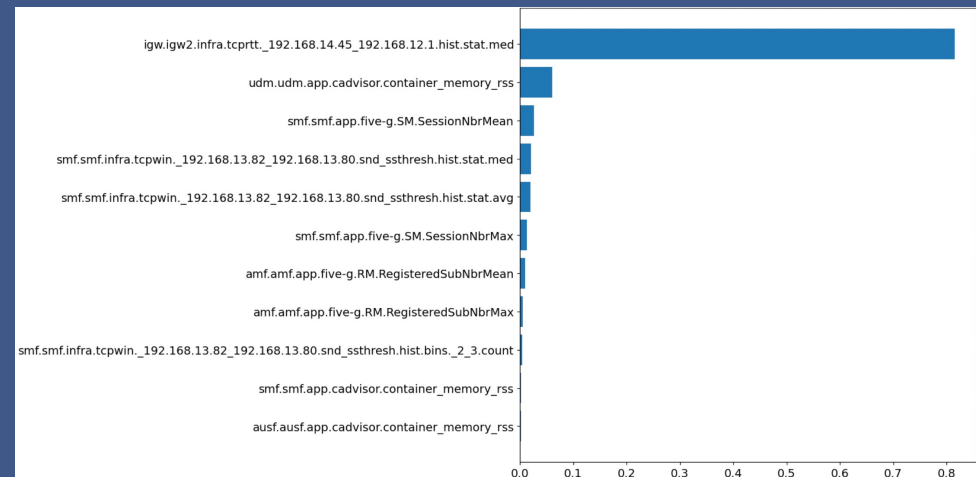

12

```
Index(['amf.amf.app.cadvisor.container_cpu_cfs_periods',
      'amf.amf.app.cadvisor.container_cpu_cfs_throttled_periods',
      'amf.amf.app.cadvisor.container_cpu_cfs_throttled_seconds',
      'amf.amf.app.cadvisor.container_cpu_system_seconds',
      'amf.amf.app.cadvisor.container_cpu_usage_seconds',
      'amf.amf.app.cadvisor.container_cpu_user_seconds',
      'amf.amf.app.cadvisor.container_last_seen',
      'amf.amf.app.cadvisor.container_memory_cache',
      'amf.amf.app.cadvisor.container_memory_failcnt',
      'amf.amf.app.cadvisor.container_memory_failures',
      ...,
      'upf.upf3.infra.tcpstates._:ffff:10.244.0.1_LAST_ACK_CLOSE.usec.avg',
      'upf.upf3.infra.tcpstates._:ffff:10.244.0.1_LAST_ACK_CLOSE.usec.max',
      'upf.upf3.infra.tcpstates._:ffff:10.244.0.1_LAST_ACK_CLOSE.usec.med',
      'upf.upf3.infra.tcpstates._:ffff:10.244.0.1_LAST_ACK_CLOSE.usec.min',
      'upf.upf3.infra.tcpstates._:ffff:10.244.0.1_SYN_RECV_ESTABLISHED.count',
      'upf.upf3.infra.tcpstates._:ffff:10.244.0.1_SYN_RECV_ESTABLISHED.usec.avg',
      'upf.upf3.infra.tcpstates._:ffff:10.244.0.1_SYN_RECV_ESTABLISHED.usec.max',
      'upf.upf3.infra.tcpstates._:ffff:10.244.0.1_SYN_RECV_ESTABLISHED.usec.med',
      'upf.upf3.infra.tcpstates._:ffff:10.244.0.1_SYN_RECV_ESTABLISHED.usec.min',
      'upf.upf3.infra.tcptracer._192.168.13.70_connect.count'],
```


Reduce metrics as much as possible from over 3000 metrics

# Method(Task2)

- method1 (difference from normal)
  - $|\text{mean(loss)}| > 2|\text{mean(normal)}|$
  - feature size: 3326 → 329
- method2 (feature importance)
  - calculate feature importance by training random forests
  - sort features in order of importance and use for LSTM training
  - Feature size: 1723, 1500, 1000, 500, 400, 329
  - Detection time: 140s



# Difference from normal model has best performance



The F1 score for the diff model was 0.97 with detection time 140s.  
The diff model has higher performance than the removed model.  
The model with 3000 features reduced had the best performance.

# Differencing process is effective

| Difference from normal |          |         |             |
|------------------------|----------|---------|-------------|
| Detection time         | F1 score | MSE     | Metrics num |
| 250                    | 1.00     | 0.00249 | 329         |
| 200                    | 1.00     | 0.00215 | 329         |
| 150                    | 0.98     | 0.00232 | 329         |
| 140                    | 0.97     | 0.00279 | 329         |
| 130                    | 0.88     | 0.00449 | 329         |
| 120                    | 0.50     | 0.00676 | 329         |

| Feature importance |          |         |             |
|--------------------|----------|---------|-------------|
| Detection time     | F1 score | MSE     | Metrics num |
| 140                | 0.93     | 0.00444 | 1723        |
| 140                | 0.88     | 0.00411 | 1500        |
| 140                | 0.89     | 0.00326 | 1000        |
| 140                | 0.75     | 0.00592 | 500         |
| 140                | 0.79     | 0.00757 | 400         |
| 140                | 0.73     | 0.00626 | 329         |

- The difference from normal model has better performance than the feature importance model under the same condition.
- The differencing process was shown to be effective.

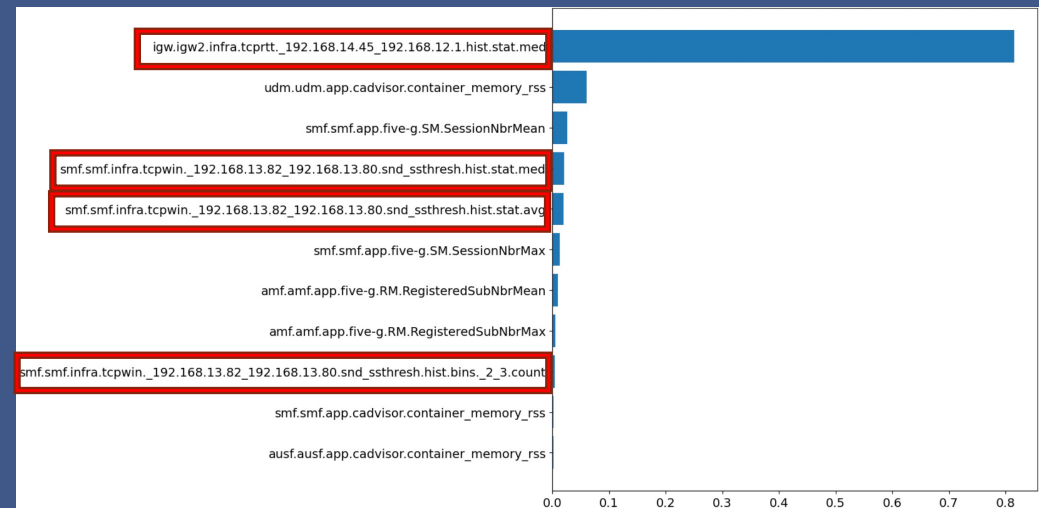
# Feature analysis

## Difference from normal model-specific metrics

```
[ 'amf.amf.infra.tcpwin._192.168.13.80_192.168.13.70.sk_wmem_queued.hist.bins._2048_4095.count',  
  'amf.amf.infra.tcpwin._192.168.13.80_192.168.13.70.sk_wmem_queued.hist.stat.med',  
  'amf.amf.infra.tcpwin._192.168.13.80_192.168.13.70.snd_cwnd.hist.bins._0_1.count',  
  'amf.amf.infra.tcpwin._192.168.13.80_192.168.13.70.snd_cwnd.hist.bins._2_3.count',  
  'amf.amf.infra.tcpwin._192.168.13.80_192.168.13.70.snd_cwnd.hist.bins._4_7.count',  
  'amf.amf.infra.tcpwin._192.168.13.80_192.168.13.70.snd_ssthresh.hist.bins._2_3.count',  
  'amf.amf.infra.tcpwin._192.168.13.80_192.168.13.72.sk_wmem_queued.hist.bins._1024_2047.count',  
  'amf.amf.infra.tcpwin._192.168.13.80_192.168.13.72.sk_wmem_queued.hist.bins._4096_8191.count',  
  'amf.amf.infra.tcpwin._192.168.13.80_192.168.13.72.snd_cwnd.hist.bins._0_1.count',  
  'amf.amf.infra.tcpwin._192.168.13.80_192.168.13.72.snd_cwnd.hist.bins._2_3.count',  
  'amf.amf.infra.tcpwin._192.168.13.80_192.168.13.72.snd_cwnd.hist.bins._4_7.count',  
  'amf.amf.infra.tcpwin._192.168.13.80_192.168.13.72.snd_ssthresh.hist.bins._2_3.count',  
  'amf.amf.infra.tcpwin._192.168.13.80_192.168.13.82.sk_wmem_queued.hist.bins._4096_8191.count',  
  'amf.amf.infra.tcpwin._192.168.13.80_192.168.13.82.snd_cwnd.hist.bins._0_1.count',  
  'amf.amf.infra.tcpwin._192.168.13.80_192.168.13.82.snd_cwnd.hist.bins._2_3.count',  
  'amf.amf.infra.tcpwin._192.168.13.80_192.168.13.82.snd_cwnd.hist.bins._4_7.count',  
  'amf.amf.infra.tcpwin._192.168.13.80_192.168.13.82.snd_ssthresh.hist.bins._2_3.count',  
  'ausf.ausf.infra.runqlat.hist.bins._131072_262143.count',  
  'ausf.ausf.infra.runqlat.hist.bins._8192_16383.count',  
  'ausf.ausf.infra.tcpcdrop._192.168.13.72_192.168.13.80_ESTABLISHED.count',  
  'ausf.ausf.infra.tcpretrans._192.168.13.72_192.168.13.70.retransmit_count',  
  'ausf.ausf.infra.tcpretrans._192.168.13.72_192.168.13.70.retransmit_rate',
```

Count metrics are important

## Feature importance

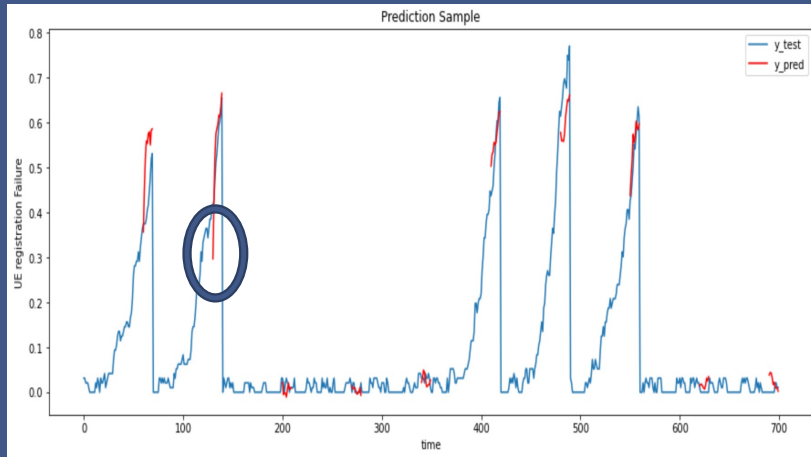


TCP metrics are important  
in loss failure




# Visualization

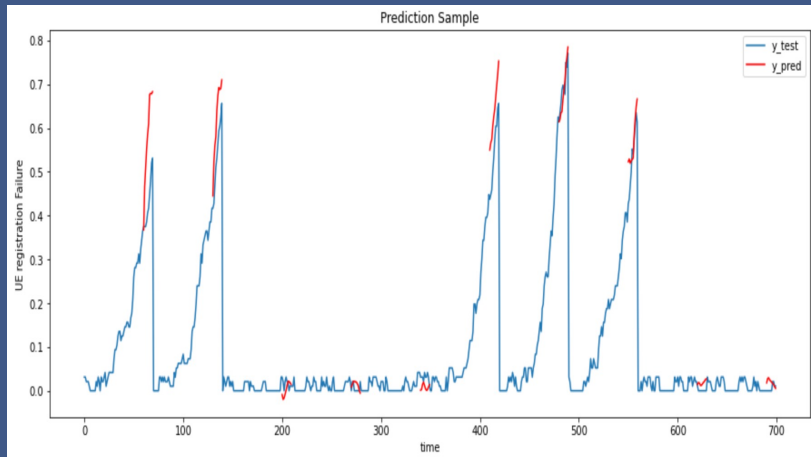
all model



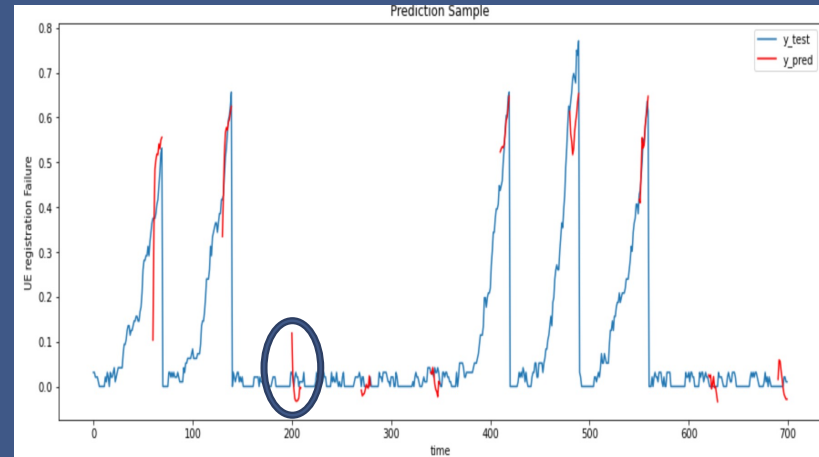
removed model



Difference from normal model



Feature importance model



# Conclusion

| Model name                   | Detection time (s) | F1 score | Metrics num |
|------------------------------|--------------------|----------|-------------|
| all model                    | 150                | 0.91     | 3325        |
| Removed model                | 140                | 0.92     | 1723        |
| Difference from normal model | 140                | 0.97     | 329         |
| Feature importance model     | 140                | 0.93     | 1723        |

- The difference from normal model is the best solution.
- Successfully predicted with a detection time of 140s
- Improved model performance by reducing the number of features
- The model performance was most improved after reducing the number of features by 3,000

# Future work

- Prediction for other failure scenarios such as CPU overload
- Predictions for different user traffic characteristics
- Analysis of metrics that are important for many factors such as failure scenarios, traffic characteristics, network size, network topology etc.
- Analysis of when the model will need to be re-trained or changed

Thank you for listening

# Appendix Task1 Table

| all            |          |         |             |
|----------------|----------|---------|-------------|
| Detection time | F1 score | MSE     | Metrics num |
| 250            | 1.00     | 0.00198 | 3325        |
| 200            | 1.00     | 0.00199 | 3325        |
| 150            | 0.91     | 0.00294 | 3325        |
| 140            | 0.89     | 0.00317 | 3325        |
| 130            | 0.80     | 0.00344 | 3325        |
| 120            | 0.44     | 0.00533 | 3325        |


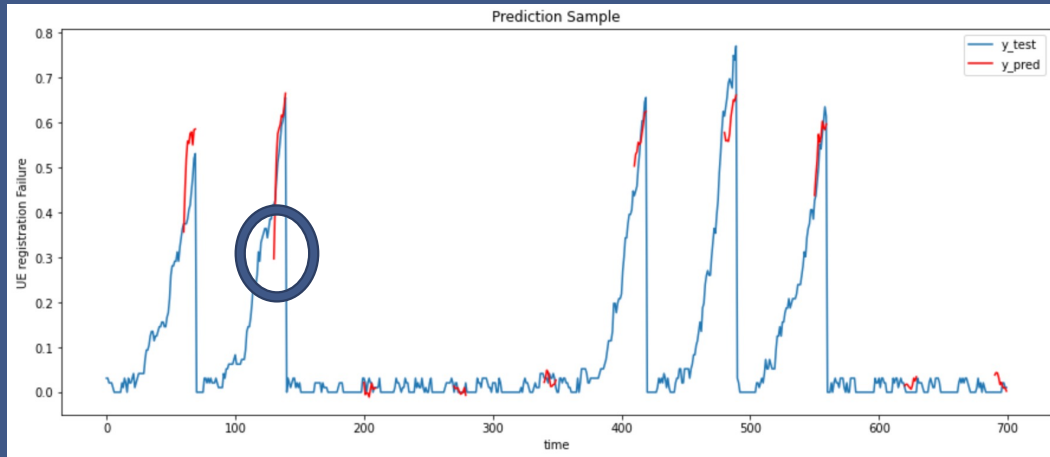
| removed |          |         |             |
|---------|----------|---------|-------------|
| removed | F1 score | MSE     | Metrics num |
| 250     | 1.00     | 0.00196 | 1723        |
| 200     | 1.00     | 0.00209 | 1723        |
| 150     | 0.98     | 0.00268 | 1723        |
| 140     | 0.92     | 0.00301 | 1723        |
| 130     | 0.70     | 0.00498 | 1723        |
| 120     | 0.39     | 0.00606 | 1723        |

Successful prediction at detection time 150s by using all features

Successful prediction at detection time 140s by removing 0 metrics

The removed model has higher performance than the all model.

# Appendix Task1 visualization



Both models may underpredict the number of UE registration failures at 600 seconds

# Appendix Task2 Table

| removed |          |         |             |
|---------|----------|---------|-------------|
| removed | F1 score | MSE     | Metrics num |
| 250     | 1.00     | 0.00196 | 1723        |
| 200     | 1.00     | 0.00209 | 1723        |
| 150     | 0.98     | 0.00268 | 1723        |
| 140     | 0.92     | 0.00301 | 1723        |
| 130     | 0.70     | 0.00498 | 1723        |
| 120     | 0.39     | 0.00606 | 1723        |


| Differenence from normal |          |         |             |
|--------------------------|----------|---------|-------------|
| diff                     | F1 score | MSE     | Metrics num |
| 250                      | 1.00     | 0.00249 | 329         |
| 200                      | 1.00     | 0.00215 | 329         |
| 150                      | 0.98     | 0.00232 | 329         |
| 140                      | 0.97     | 0.00279 | 329         |
| 130                      | 0.88     | 0.00449 | 329         |
| 120                      | 0.50     | 0.00676 | 329         |

The F1 score for the diff model was 0.97 with detection time 140s.


The diff model has higher performance than the removed model.

The model with 3000 features reduced had the best performance.

# Appendix Task2 visualization





The removed model may underpredict the number of UE registration failures at 600 seconds.



The diff model may overpredict.



# Appendix task2 visualization



The diff model predicts more accurately than the RF model.

The RF model does not accurately predict normal phase.